

CLaFICLe: Cross Lingual Adaptation for In-Context Learning

Giulio Starace

University of Amsterdam / Amsterdam, The Netherlands

giulio.starace@gmail.com

Abstract

This document is a supplement to the general instructions for *ACL authors. It contains instructions for using the \LaTeX style files for ACL conferences. The document itself conforms to its own specifications, and is therefore an example of what your manuscript should look like. These instructions should be used both for papers submitted for review and for final versions of accepted papers.

1 Introduction

2 Related Work

3 Method

textwidth in inches: 6.30045in

columnwidth in inches: 3.03209in

4 Results and Discussion

5 Conclusion

References

Sewon Min, Mike Lewis, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2022. [MetaICL: Learning to Learn In Context](#).

Benjamin Minixhofer, Fabian Paischer, and Navid Rekasbas. 2022. [WECHSEL: Effective initialization of subword embeddings for cross-lingual transfer of monolingual language models](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 3992–4006, Seattle, United States. Association for Computational Linguistics.

6 Appendices

Use `\appendix` before any appendix section to switch the section numbering over to letters. See Appendix 7 for an example.

7 Example Appendix

This is an appendix.

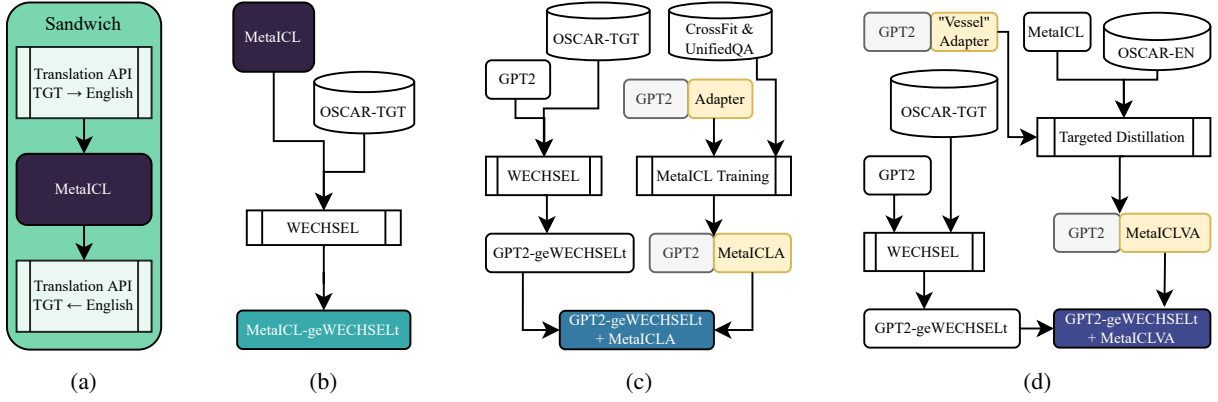


Figure 1: Overview of each of the models evaluated in one of the two TGT languages (French or German). The baseline **sandwich** model (a) sandwiches **MetaICL** (Min et al., 2022) (which we separately evaluate only in English) between two complementary translation API calls. **MetaICL-geWECHSELt** (b) is the result of applying **WECHSEL** (Minixhofer et al., 2022) to **MetaICL**. **GPT2-geWECHSELt+MetaICLA** combines **MetaICLA**, an adapter trained on the **MetaICL** dataset and objective, with a TGT-language GPT2 base obtained via **WECHSEL**. **GPT2-geWECHSELt+MetaICLVA** does the same, except **MetaICLVA** is trained via targeted distillation with supervision provided by **MetaICL**. For more details, refer to section 3.

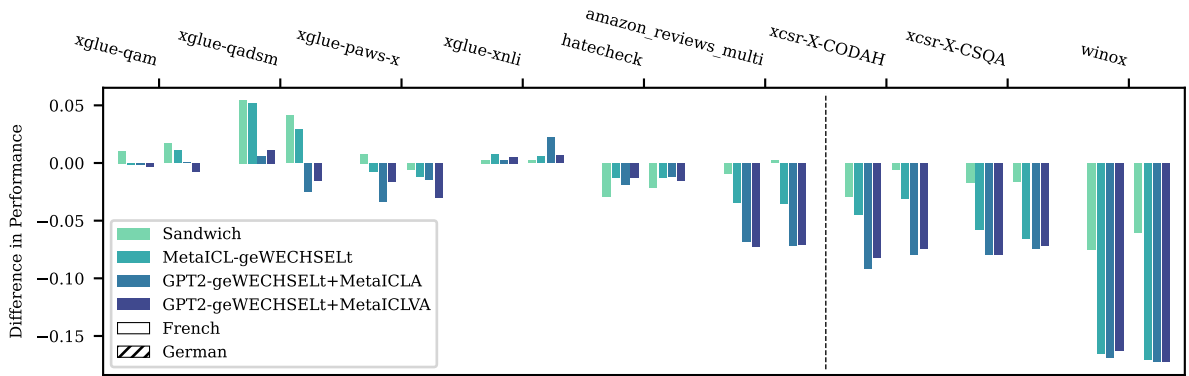


Figure 2: Performance gap on our multi-task benchmark between each of the language-adapted models in our target languages and **MetaICL** in English. Positive values indicate that the adapted models are outperforming **MetaICL**, while negative values indicate the reverse. The dashed line separates whether a given task uses accuracy (left) or F1-score (right) as the performance metric.